

---

# National University of Computer and Emerging Sciences

Spring 2023

Inter-Process Communication

---

# Interprocess Communication

- A process has access to the memory which constitutes its own address space.
- When a child process is created, the only way to communicate between a parent and a child process is:
  - The variables are replicas
  - The parent receives the exit status of the child
- So far, we've discussed communication mechanisms only during process creation/termination
- Processes may need to communicate during their life time.

# Cooperating Processes

- *Independent* process cannot affect or be affected by the execution of another process.
- *Cooperating* process can affect or be affected by the execution of another process
- Advantages of process cooperation
  - Information sharing
  - Computation speed-up:
    - make use of multiple processing elements
  - Modularity
  - Convenience:
    - editing, printing, compiling in parallel
- Dangers of process cooperation
  - Data corruption, deadlocks, increased complexity
  - Requires processes to synchronize their processing

# Purposes for IPC

- IPC allows processes to communicate and synchronize their actions without sharing the same address space
  - ❑ Data Transfer
  - ❑ Sharing Data
  - ❑ Event notification
  - ❑ Resource Sharing and Synchronization

# IPC Mechanisms

- Mechanisms used for communication and synchronization
  - Message Passing
    - message passing interfaces, mailboxes and message queues
    - sockets, pipes
  - *Shared Memory*: Non-message passing systems
  - *Synchronization* – primitives such as semaphores to higher level mechanisms such as monitors
  - *Event Notification* - UNIX *signals*
- We will defer a detailed discussion of synchronization mechanisms and concurrency until a later class
- Here we want to focus on some common (and fundamental) IPC mechanisms

# Message Passing

- In a *Message system* there are no shared variables.
- IPC facility provides two operations for fixed or variable sized message:
  - *send(message)*
  - *receive(message)*
- If processes *P* and *Q* wish to communicate, they need to:
  - establish a *communication link*
  - exchange messages via *send and receive*
- Implementation of communication link
  - physical (e.g., memory, network etc)
  - logical (e.g., syntax and semantics, abstractions)

# Message Passing Systems

- Exchange messages over a communication link
- Methods for implementing the communication link and primitives (*send/receive*):
  1. Direct or Indirect communications (*Naming*)
  2. Symmetric or Asymmetric communications (blocking versus non-blocking)
  3. Buffering
  4. Send-by-copy or send-by-reference
  5. fixed or variable sized messages

# Direct Communication – Internet and Sockets

- Processes must name each other *explicitly*:
  - Symmetric Addressing
    - *send(P, message)* – send to process P
    - *receive(Q, message)* – receive from Q
  - Asymmetric Addressing
    - *send(P, message)* – send to process P
    - *receive(id, message)* – rx from any; system sets id = sender
- Properties of communication link
  - Links established automatically between pairs
  - processes must know each others ID
  - Exactly one link per pair of communicating processes
- Disadvantage: a process must know the name or ID of the process(es) it wishes to communicate with



# Indirect Communication

- Messages are sent to or received from *mailboxes* (also referred to as *ports*).
  - Each mailbox has a unique id.
  - Processes can communicate only if they share a mailbox.
- Primitives:
  - *send(A, message)* – send a message to mailbox A
  - *receive(A, message)* – receive a message from mailbox A
- Properties of communication link
  - Link established only if processes share a common mailbox
  - A link may be associated with more than 2 processes.
  - Each pair of processes may share several communication links.

---

# Indirect Communication- Ownership

- process owns (i.e. mailbox is implemented in user space):
    - ❑ only the owner may receive messages through this mailbox.
    - ❑ Other processes may only send.
    - ❑ When process terminates any “owned” mailboxes are destroyed.
  - kernel owns
    - ❑ then mechanisms provided to create, delete, send and receive through mailboxes.
    - ❑ Process that creates mailbox owns it (and so may receive through it)
    - ❑ but may transfer ownership to another process.
-

# Indirect Communication

- Mailbox sharing:
  - $P_1$ ,  $P_2$ , and  $P_3$  share mailbox  $A$ .
  - $P_1$  sends;  $P_2$  and  $P_3$  receive.
  - Who gets the message?
- Solutions
  - Allow a link to be associated with at most two processes.
  - OR Allow only one process at a time to execute a receive operation.
  - OR Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was

# Synchronization

- Message passing may be either *blocking* or *non-blocking*.
  - blocking send:
    - sender blocked until message received by mailbox or process
  - nonblocking send:
    - sender resumes operation immediately after sending
  - blocking receive:
    - receiver blocks until a message is available
  - nonblocking receive:
    - receiver returns immediately with either a valid or null message.

# Buffering

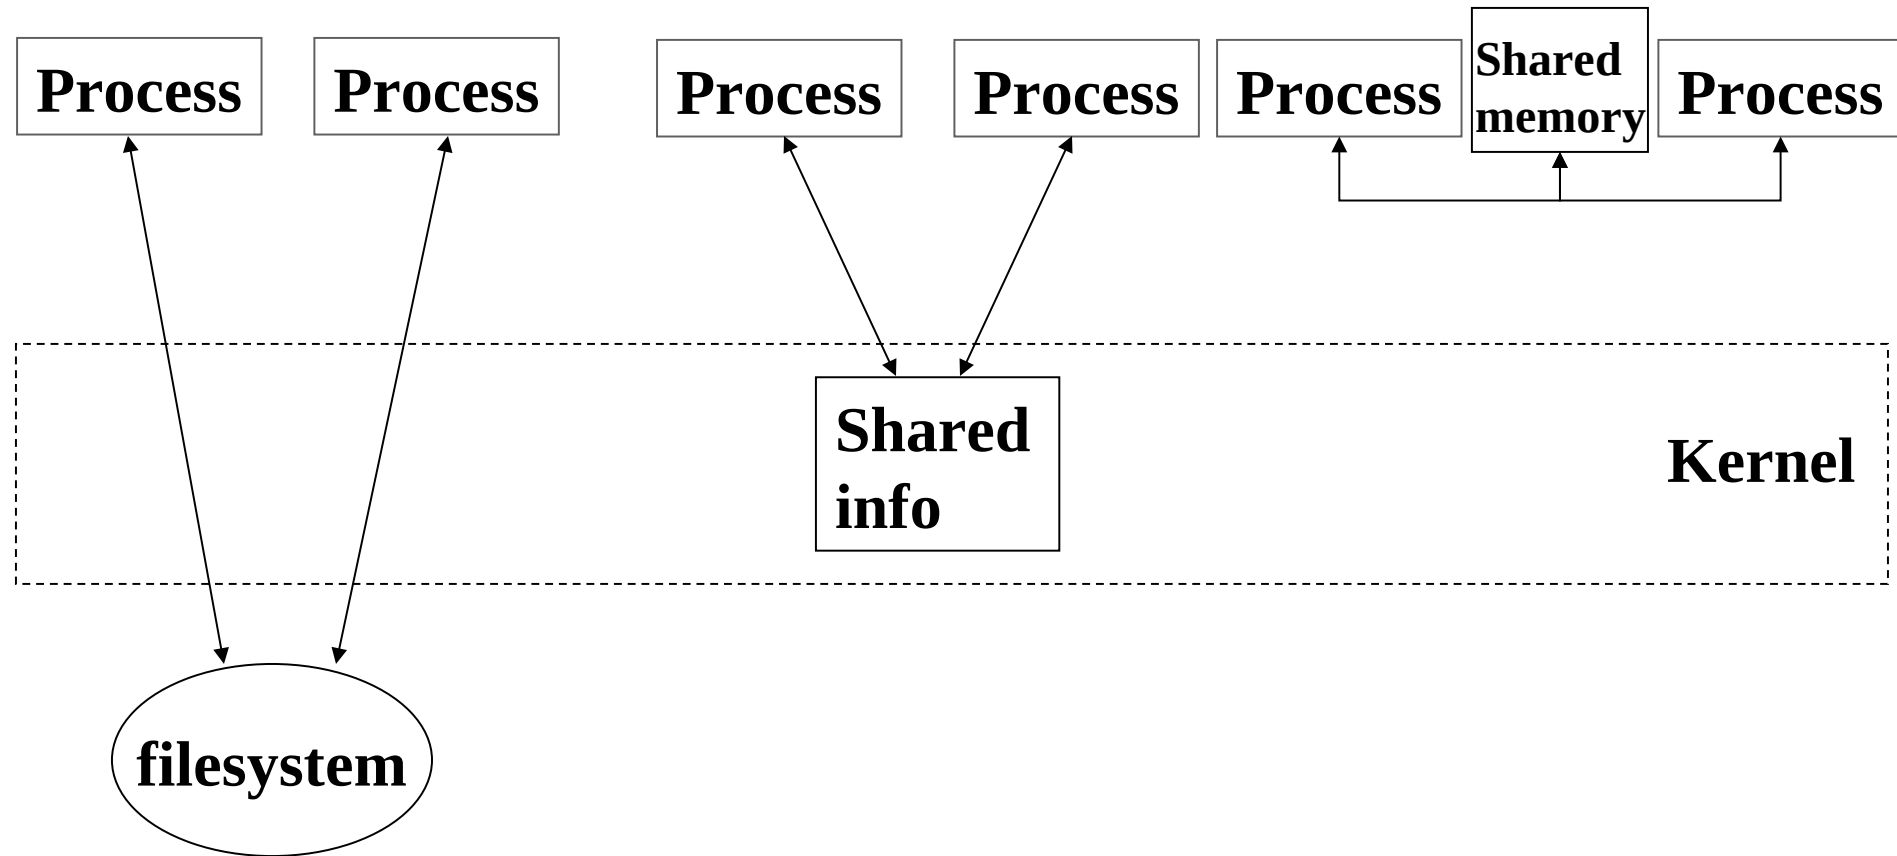
- All messaging system require framework to temporarily buffer messages.
- These queues are implemented in one of three ways:
  1. **Zero capacity**

No messages may be queued within the link, requires sender to block until receiver retrieves message.
  2. **Bounded capacity**

Link has finite number of message buffers. If no buffers are available then sender must block until one is freed up.
  3. **Unbounded capacity**

Link has unlimited buffer space, consequently send never needs to block.

# Unix IPC




# Unix IPC

- Message passing
  - Pipes
  - FIFOs
  - Message queues
- Synchronization
  - Mutexes
  - Condition variables
  - read-write locks
  - Semaphores
- Shared memory
  - Anonymous
  - Named
- Procedure calls
  - RPC
- Event notification
  - Signals

# File Descriptors

- The PCB (*task\_struct*) of each process contains a pointer to a *file\_struct*





# File Descriptors

- The *files\_struct* contains pointers to file data structures
- Each one describes a file being used by this process.
- *f\_mode*: describes file mode, read only, read and write or write only.
- *f\_pos*: holds the position in the file where the next read or write operation will occur.
- *f\_inode*: points at the actual file

# File Descriptors

- Every time a file is opened, one of the free file pointers in the *files\_struct* is used to point to the new file structure.
- Linux processes expect three file descriptors to be open when they start.
- These are known as *standard input*, *standard output* and *standard error*


# File Descriptors

- The program treat them all as files.
- These three are usually inherited from the creating parent process.
- All accesses to files are via standard system calls which pass or return file descriptors.
- *standard input, standard output and standard error* have file descriptors 0, 1 and 2.

# File Descriptors

- `char buffer[10];`
- Read from standard input (by default it is keyboard)
  - `read(0,buffer,5);`
- Write to standard output (by default is is monitor))
  - `write(1,buffer,5);`
- By changing the file descriptors we can write to files
- `fread/fwrite` etc are wrappers around the above `read/write` functions

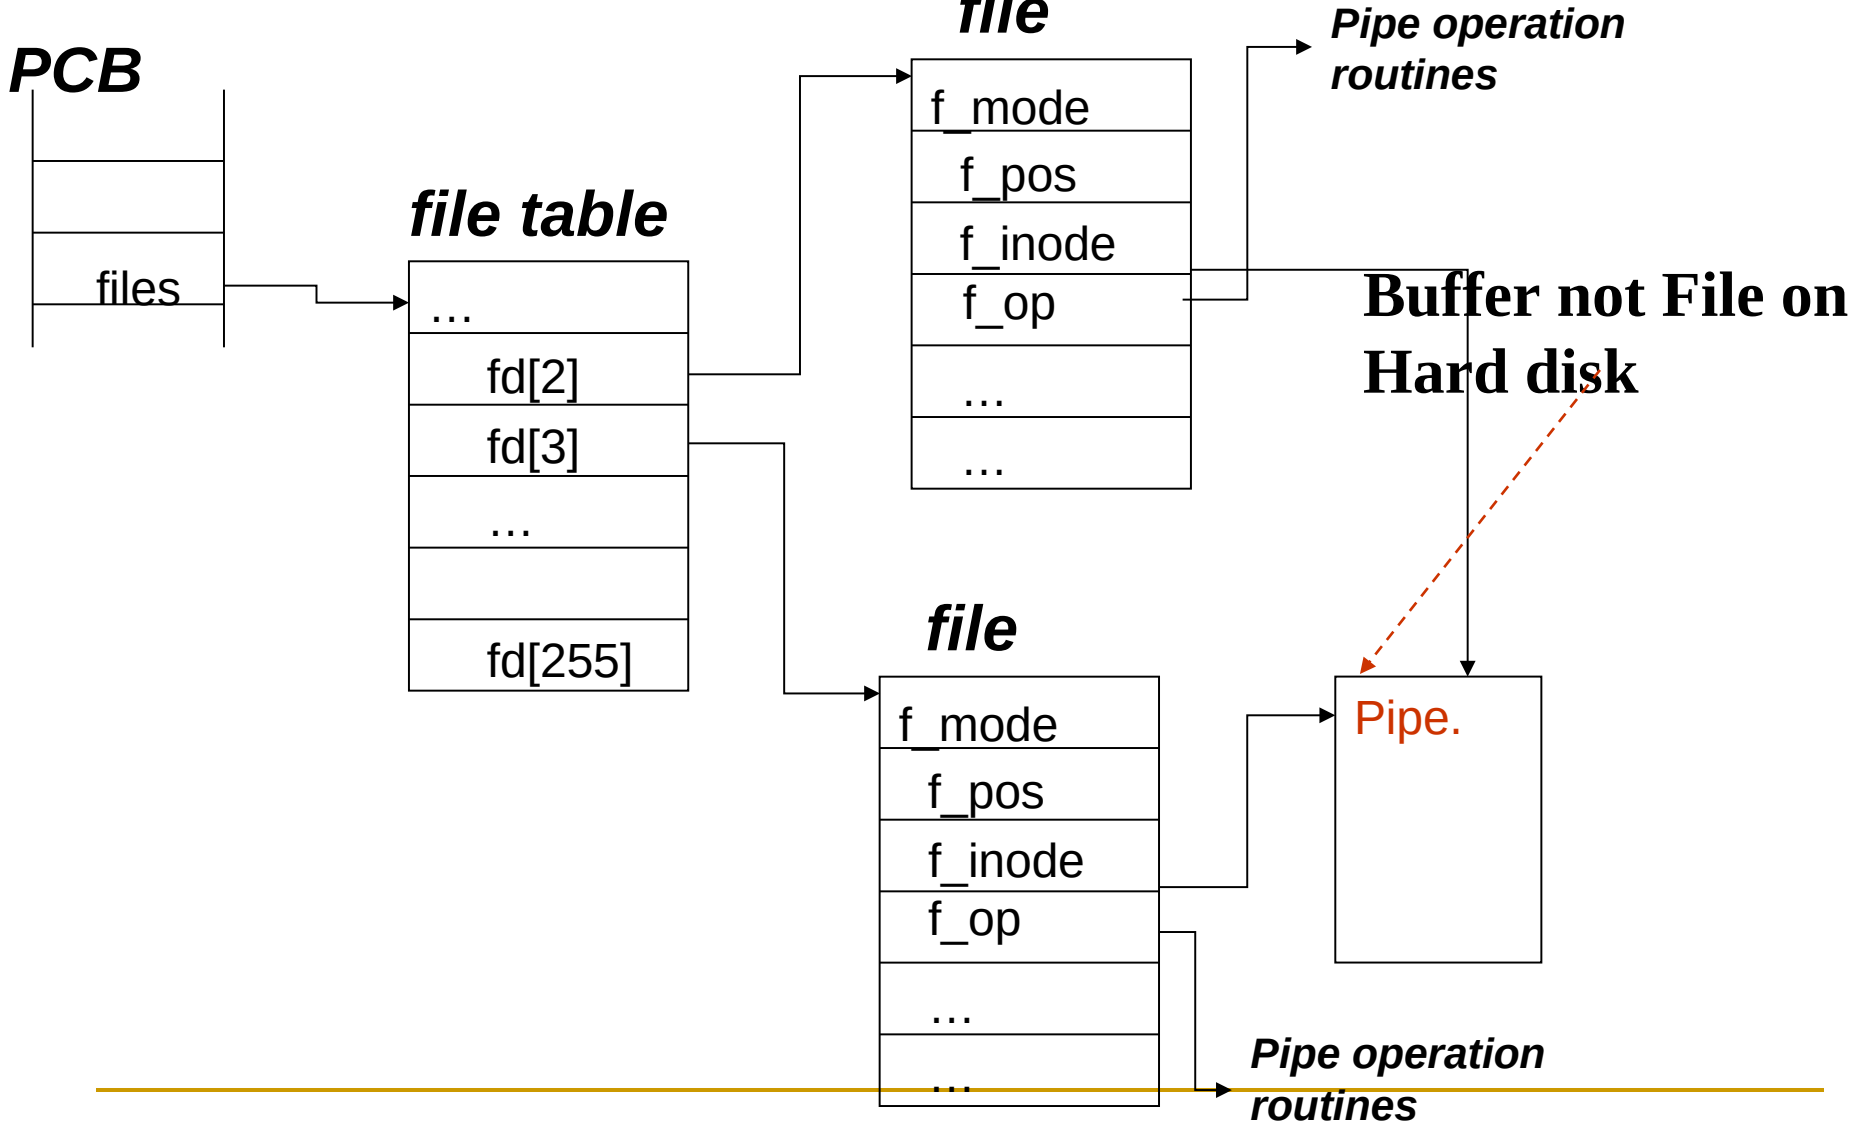
# Pipes: Shared info in kernel's memory



# Pipes

- A pipe is implemented using two file data structures which both point at the same temporary data node.
- This hides the underlying differences from the generic system calls which read and write to ordinary files
- Thus, reading/writing to a pipe is similar to reading/writing to a file

# Pipes



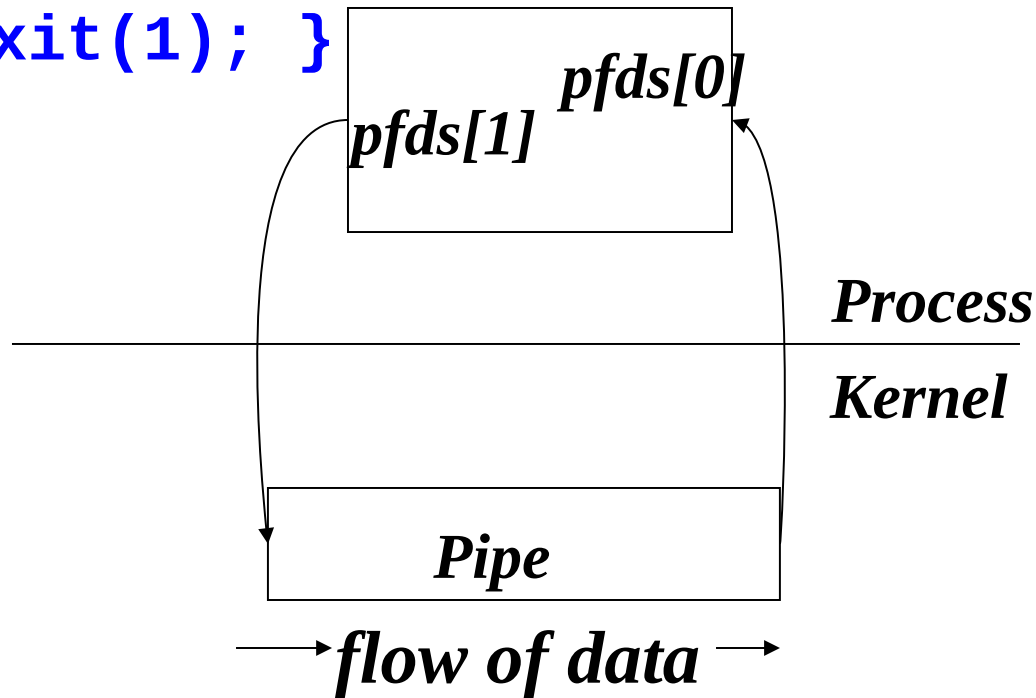
# Pipe Creation

- `#include <unistd.h>`
- `int pipe(int fildes[2]);`
- Creates a pair of file descriptors pointing to a pipe inode
- Places them in the array pointed to by *fildes*
- *fildes[0]* is for reading
- *fildes[1]* is for writing.
- On success, zero is returned.
- On error, -1 is returned



# Pipe Creation

```
int main()  
{ int pfd[2];  
  if (pipe(pfd) == -1)  
    { perror("pipe");  
      exit(1); }  
}
```



# Reading/Writing from/to a Pipe

- `int read(int filedescriptor, char *buffer, int bytetoread);`
- `int write(int filedescriptor, char *buffer, int bytetowrite);`

# Example

```
int main()
{ int pfd[2];
  char buf[30];
  if (pipe(pfd) == -1) {
    perror("pipe");
    exit(1); }
  printf("writing to file descriptor %d\n",
    pfd[1]);
  write(pfd[1], "test", 5);
  printf("reading from file descriptor %d\n", pfd[0]);
  read(pfd[0], buf, 5);
  read(0, buf, 5);
}
```


# A Channel between two processes

- Remember: the two processes have a parent / child relationship
- The child was created by a `fork()` call that was executed by the parent.
- The child process is an image of the parent process
- Thus, all the **file descriptors** that are opened by the parent are now available in the child.

# A Channel between two processes

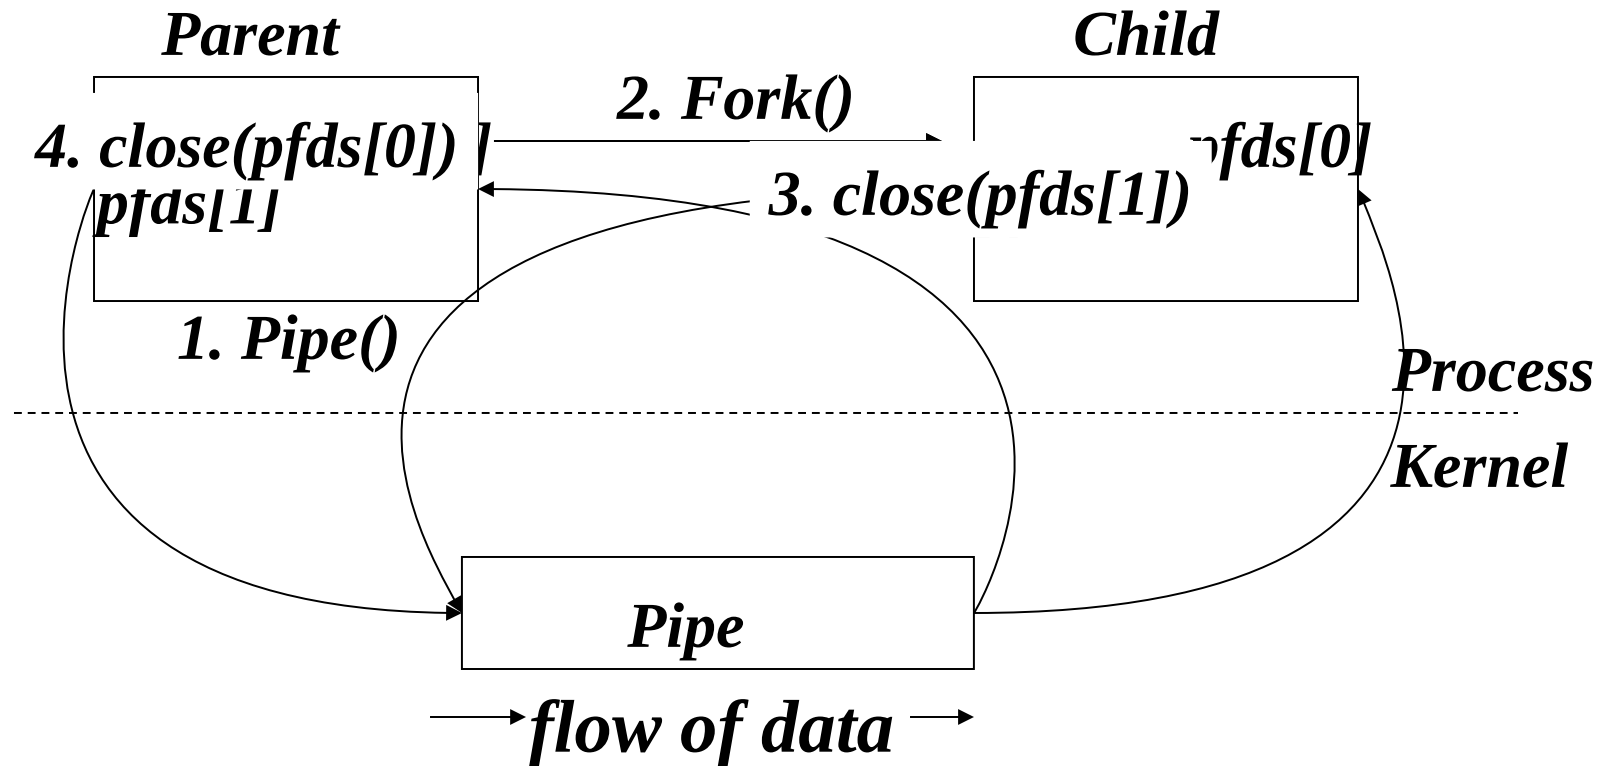
- The **file descriptors** refer to the same I/O entity, in this case a pipe.
- The pipe is inherited by the child
- And may be passed on to the grand-children by the child process or other children by the parent.

# A Channel between two processes



# A Channel between two processes

- To allow one way communication each process should **close** one end of the pipe.



# Closing the pipe

- The **file descriptors** associated with a pipe can be closed with the `close(fd)` system call



# An Example of pipes with

## fork

```
int main() {
    int pfd[2];
    char buf[30];
    pipe(pfd); .....1
    if (!fork()) .....2
    { close(pfd[0]); .....3
        printf(" CHILD: writing to the pipe\n");
        write(pfd[1], "test", 5);
        printf(" CHILD: exiting\n");
        exit(0);
    }
    else { close(pfd[1]); .....4
        printf("PARENT: reading from pipe\n");
        read(pfd[0], buf, 5);
        printf("PARENT: read \"%s\"\n", buf);
        wait(NULL);
    }
}
```